

PROJECT SPECIFICATION

Create Your Own Image Classifier

Files Submitted

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
Submission Files	The submission includes all required files. (Model checkpoints not required.)	<p>The required files are submitted as part of the repo (root directory). These include</p> <ul style="list-style-type: none">• Image_Classifier_Project.ipynb• train.py• predict.py

Part 1 - Development Notebook

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
Package Imports	All the necessary packages and modules are imported in the first cell of the notebook	The imports are done in cell <1>
Training data augmentation	torchvision transforms are used to augment the training data with random scaling, rotations, mirroring, and/or cropping	<p>The training and test data is augmented using pytorch transform functions.</p> <p>From training data the transforms used are:</p> <p>RandomRotation(),</p> <p>RandomResizedCrop() and</p> <p>RandomHorizontalFlip().</p> <p>For test and validation, the transforms used are:</p> <p>Resize()</p> <p>CenterCrop().</p> <p>I tried additional techniques but the performance suffered. I tried the same settings for test and validation as training using flips and random crop!</p> <p>Again the network was not yielding the accuracy of more than 80% and test accuracy is always below 80%. Even if I tested with 50 epochs, the network</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>was struggling.</p> <p>This can be observed in cell <3></p>
Data normalization	The training, validation, and testing data is appropriately cropped and normalized	<p>The normalization is done using torchvision Normalize() method for all three (training, test and validation) data sets</p> <p>This can be observed in cell <3></p>
Data batching	The data for each set is loaded with torchvision's DataLoader	<p>The data is loaded by creating a DataLoader object (trainloader, testloader, validloader) for each of the data sets (train, test and validation) respectively.</p> <p>This can be observed in cell <3></p>
Data loading	The data for each set (train, validation, test) is loaded with torchvision's ImageFolder	<p>The data for each set is loaded with torchvision's ImageFolder method.</p> <p>This can be observed in cell <3></p>
Pretrained Network	A pretrained network such as VGG16 is loaded from torchvision.models and the parameters are frozen	<p>The model is initialized using vgg16 network.</p> <p>This can be seen in cell<></p> <p>The parameters are frozen using</p> <p>params.requires_grad = False</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		This can be observed in cell <11>
Feedforward Classifier	A new feedforward network is defined for use as a classifier using the features as input	<p>A new classifier is created using</p> <pre> classifier = nn.Sequential(nn.Linear(25088, 4096), nn.ReLU(), nn.Dropout(p=0.2), nn.Linear(4096, 102), nn.LogSoftmax(dim=1)) </pre> <p>and the model is updated with</p> <pre> model.classifier = classifier. </pre> <p>This can be observed in cell <5> as part of the build_model function</p>
Training the network	The parameters of the feedforward classifier are appropriately trained, while the parameters of the feature network are left	<p>The first few times the model is trained using the code in cells <></p> <p>Later on, the code is converted into functions and the retrained for better performace.</p> <p>The final network parameters I have chosen given me close to 89%</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
	static	<p>validation accuracy and the test accuracy is around 85%</p> <p>the poor performing results can be observed in cells <10(vgg16), 21, 26, 10(vgg19)></p> <p>More over in the notebook, Image_Classification_Project-experiments.ipynb, the other training I have done is preserved. But I cleaned up the main notebook for submission.</p>
Testing Accuracy	The network's accuracy is measured on the test data	The final test accuracy is <cell 10(vgg16)>
Validation Loss and Accuracy	During training, the validation loss and accuracy are displayed	<p>During training I displayed the following info. Every time I run something, i'm a fan of estimating how much each step/epoch takes. So the information I captured are:</p> <p>Epochs</p> <p>Steps [time taken for steps]</p> <p>Training Loss</p> <p>Validation Loss</p> <p>Validation accuracy</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>The format is as below</p> <p>Epoch: 1/5 Steps: 30[14.202s] Training Loss: 7.513 Validation Loss: 3.203 Validation Accuracy: 0.338</p>
Loading checkpoints	There is a function that successfully loads a checkpoint and rebuilds the model	A function (load_model()) is defined to load the checkpoint. This can be seen in cell <41>
Saving the model	The trained model is saved as a checkpoint along with associated hyperparameters and the class_to_idx dictionary	A function (savemodel()) is defined to save the model parameters. This can be observed in cell 32
Image Processing	The process_image function successfully converts a PIL image into an object that can be used as input to a trained model	process_image() function is defined to address this rubric in cell<47>
Class Prediction	The predict function	Predict() function is defined to address this rubric in cell <49>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
	successfully takes the path to an image and a checkpoint, then returns the top K most probably classes for that image	
Sanity Checking with matplotlib	A matplotlib figure is created displaying an image and its associated top 5 most probable classes with actual flower names	Sanity check is done in cell <67> a function sanity_check() is subsequently defined to do the same task.

Part 2 - Command Line Application

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
Training a network	train.py successfully trains a new network on a dataset of images and saves the model to a checkpoint	<p>Please take a look at the output here.</p> <p>(dsnd-pytorch)</p> <pre>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ python train.py --arch 'vgg16' --n_inputs 25088 --n_hidden 4096 --n_classes 102 --epochs 1 --learning_rate 0.001 --batch_size 32 --drop_prob 0.5 --gpu -- print_every 30 --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth</pre> <p>Using cuda:0</p> <p>THCudaCheck FAIL file=/opt/conda/conda-bld/pytorch_1544176307774/work/ aten/src/THC/THCGeneral.cpp line=405 error=11 : invalid argument</p> <p>Epoch: 1/1 Steps: 30[14.464s] Training Loss: 6.324 Validation Loss: 2.946 Validation Accuracy: 0.390</p> <p>Epoch: 1/1 Steps: 60[29.152s] Training Loss: 3.083 Validation Loss: 2.122 Validation Accuracy: 0.497</p> <p>Epoch: 1/1 Steps: 90[43.844s] Training Loss: 2.749 Validation Loss: 1.579 Validation Accuracy: 0.608</p> <p>Epoch: 1/1 Steps: 120[58.368s] Training Loss: 2.375 Validation Loss:</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>1.429 Validation Accuracy: 0.640</p> <p>Epoch: 1/1 Steps: 150[72.662s] Training Loss: 2.212 Validation Loss: 1.340 Validation Accuracy: 0.634</p> <p>Epoch: 1/1 Steps: 180[86.997s] Training Loss: 2.106 Validation Loss: 1.121 Validation Accuracy: 0.699</p> <p>Accuracy achieved by the network on test images is: 73%</p> <pre> gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ ls -l vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth -rw-r--r-- 1 gvenkat gvenkat 1297067396 Jan 31 05:44 vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth (dsnd-pytorch) gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ </pre>
Training validation log	The training loss, validation loss, and validation accuracy are printed out as a network trains	<p>From the previous run we clearly see that the training loss, validation loss and validation accuracy are all reported for each run.</p> <pre> Epoch: 1/1 Steps: 30[14.464s] Training Loss: 6.324 Validation Loss: 2.946 Validation Accuracy: 0.390 Epoch: 1/1 Steps: 60[29.152s] Training Loss: 3.083 Validation Loss: </pre>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>2.122 Validation Accuracy: 0.497</p> <p>Epoch: 1/1 Steps: 90[43.844s] Training Loss: 2.749 Validation Loss: 1.579 Validation Accuracy: 0.608</p> <p>Epoch: 1/1 Steps: 120[58.368s] Training Loss: 2.375 Validation Loss: 1.429 Validation Accuracy: 0.640</p> <p>Epoch: 1/1 Steps: 150[72.662s] Training Loss: 2.212 Validation Loss: 1.340 Validation Accuracy: 0.634</p> <p>Epoch: 1/1 Steps: 180[86.997s] Training Loss: 2.106 Validation Loss: 1.121 Validation Accuracy: 0.699</p>
Model architecture	The training script allows users to choose from at least two different architectures available from torchvision.models	<p>I have given an option to choose more than one model from torchvision.models. I tested with vgg16 and vgg19. And also with densenet121</p> <p>(dsnd-pytorch)</p> <pre>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ python train.py --arch 'densenet121' --n_inputs 1024 --n_hidden 120 -- n_classes 102 --epochs 1 --learning_rate 0.001 --batch_size 32 --drop_prob 0.5 --gpu --print_every 30 --checkpoint densene121_checkpoint_l1_001_bs_32_dp_05_e1.pth /home/gvenkat/anaconda3/envs/dsnd-pytorch/lib/python3.7/site-packages/</pre>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>torchvision/models/densenet.py:212: UserWarning: nn.init.kaiming_normal is now deprecated in favor of nn.init.kaiming_normal_.</p> <pre>nn.init.kaiming_normal(m.weight.data)</pre> <p>Using cuda:0</p> <p>THCudaCheck FAIL file=/opt/conda/conda-bld/pytorch_1544176307774/work/aten/src/THC/THCGeneral.cpp line=405 error=11 : invalid argument</p> <p>Epoch: 1/1 Steps: 30[12.756s] Training Loss: 4.563 Validation Loss: 4.411 Validation Accuracy: 0.138</p> <p>Epoch: 1/1 Steps: 60[25.512s] Training Loss: 4.381 Validation Loss: 4.173 Validation Accuracy: 0.160</p> <p>Epoch: 1/1 Steps: 90[38.551s] Training Loss: 4.198 Validation Loss: 3.892 Validation Accuracy: 0.213</p> <p>Epoch: 1/1 Steps: 120[51.506s] Training Loss: 4.022 Validation Loss: 3.591 Validation Accuracy: 0.237</p> <p>Epoch: 1/1 Steps: 150[64.518s] Training Loss: 3.778 Validation Loss: 3.348 Validation Accuracy: 0.261</p> <p>Epoch: 1/1 Steps: 180[78.390s] Training Loss: 3.559 Validation Loss:</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>3.092 Validation Accuracy: 0.344</p> <p>Accuracy achieved by the network on test images is: 35%</p> <p>In addition if the inputs and hidden layers are given correctly other models should also work!</p>
Model hyperparameters	The training script allows users to set hyperparameters for learning rate, number of hidden units, and training epochs	<p>The following is an example run.</p> <pre>python train.py --arch 'vgg16' --n_inputs 25088 --n_hidden 4096 --n_classes 102 --epochs 1 --learning_rate 0.001 --batch_size 32 --drop_prob 0.5 --gpu --print_every 30 --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth</pre> <p>Please see the highlighted ones.</p>
Training with GPU	The training script allows users to choose training the model on a GPU	<pre>python train.py --arch 'vgg16' --n_inputs 25088 --n_hidden 4096 --n_classes 102 --epochs 1 --learning_rate 0.001 --batch_size 32 --drop_prob 0.5 --gpu --print_every 30 --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth</pre> <p>From the output we can clearly see the script saying it's using cuda:0</p> <p>(dsnd-pytorch)</p> <pre>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$</pre>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>python train.py --arch 'vgg16' --n_inputs 25088 --n_hidden 4096 --n_classes 102 --epochs 1 --learning_rate 0.001 --batch_size 32 --drop_prob 0.5 --gpu --print_every 30 --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e1.pth</p> <p>Using cuda:0</p> <p>THCudaCheck FAIL file=/opt/conda/conda-bld/pytorch_1544176307774/work/aten/src/THC/THCGeneral.cpp line=405 error=11 : invalid argument</p> <p>Epoch: 1/1 Steps: 30[14.464s] Training Loss: 6.324 Validation Loss: 2.946 Validation Accuracy: 0.390</p> <p>Epoch: 1/1 Steps: 60[29.152s] Training Loss: 3.083 Validation Loss: 2.122 Validation Accuracy: 0.497</p> <p>Epoch: 1/1 Steps: 90[43.844s] Training Loss: 2.749 Validation Loss: 1.579 Validation Accuracy: 0.608</p> <p>Epoch: 1/1 Steps: 120[58.368s] Training Loss: 2.375 Validation Loss: 1.429 Validation Accuracy: 0.640</p> <p>Epoch: 1/1 Steps: 150[72.662s] Training Loss: 2.212 Validation Loss: 1.340 Validation Accuracy: 0.634</p> <p>Epoch: 1/1 Steps: 180[86.997s] Training Loss: 2.106 Validation Loss: 1.121 Validation Accuracy: 0.699</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<p>Accuracy achieved by the network on test images is: 73%</p> <p>(dsnd-pytorch)</p> <p>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$</p>
Predicting classes	The predict.py script successfully reads in an image and a checkpoint then prints the most likely image class and it's associated probability	<p>The following is the predict run.</p> <p>In this particular case I used top_k = 1</p> <p>(dsnd-pytorch)</p> <p>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$</p> <p>python predict.py --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e10.pth --image "flower_data/test/1/image_06754.jpg" --top_k 1 --category_names "cat_to_name.json" --gpu</p> <p>Using cuda:0</p> <p>[0.780955970287323]</p> <p>['pink primrose']</p> <p>[1]</p>
Top K classes	The predict.py script allows users to print out the top K	Please take a look at the output below. I used topk=5 and we see their

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
	classes along with associated probabilities	<p>classes and associated probabilities</p> <p>(dsnd-pytorch)</p> <pre>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ python predict.py --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e10.pth --image "flower_data/test/1/ image_06754.jpg" --top_k 5 --category_names "cat_to_name.json" --gpu</pre> <p>[0.780955970287323, 0.11035141348838806, 0.03172770142555237, 0.012553058564662933, 0.01197933778166771]</p> <p>['pink primrose', 'mexican petunia', 'petunia', 'carnation', 'mallow']</p> <p>[1, 98, 51, 31, 97]</p>
Displaying class names	The predict.py script allows users to load a JSON file that maps the class values to other category names	<p>This is implemented in predict.py in lines 172-173</p> <p>with open(args.category_names, 'r') as f:</p> <pre>cat_to_name = json.load(f)</pre>
Predicting with GPU	The predict.py script allows users to use the GPU to calculate the predictions	<p>I have given option to use gpu. From the following run, we can clearly see the same.</p> <p>(dsnd-pytorch)</p>

CRITERIA	MEETS SPECIFICATIONS	COMMENTS
		<pre>gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$ python predict.py --checkpoint vgg16_checkpoint_l1_001_bs_32_dp_05_e10.pth --image "flower_data/test/1/ image_06754.jpg" --top_k 1 --category_names "cat_to_name.json" --gpu</pre> <p>Using cuda:0</p> <pre>[0.780955970287323] ['pink primrose'] [1] (dsnd-pytorch) gvenkat@gvAwAuR7:~/DSND/DSND_Term1/projects/p2_image_classifier\$</pre>